

Introduction to

MongoDB




The Great Divide




Easy Flexible

Scalable




What is MongoDB ?

- Scalable High-Performance Open-source, Document-orientated database.
- Built for Speed
- Rich Document based queries for Easy readability.
- Full Index Support for High Performance.
- Replication and Failover for High Availability.
- Auto Sharding for Easy Scalability.
- Map / Reduce for Aggregation.



Why use MongoDB?

- SQL was invented in the 70's to store data.
- MongoDB stores documents (or) objects.
- Now-a-days, everyone works with objects (Python/Ruby/Java/etc.)
- And we need Databases to persist our objects.
Then why not store objects directly ?
- Embedded documents and arrays reduce need for joins. No Joins and No-multi document transactions.



What is MongoDB great for?

- RDBMS replacement for Web Applications.


Semi-structured Content Management.

Real-time Analytics & High-Speed Logging.

- Caching and High Scalability

Web 2.0, Media, SAAS, Gaming

HealthCare, Finance, Telecom, Government



Not great for?

- Highly Transactional Applications.
- Problems requiring SQL.


Some Companies using MongoDB in Production



Let's Dive in !




When I say



- Made up of Multiple Collections.
- Created on-the-fly when referenced for the first time.

When I say




- Schema-less, and contains Documents.

Indexable by one/more keys.

- Created on-the-fly when referenced for the first time.

Capped Collections: Fixed size, older records get dropped after reaching the limit.


When I say



- Stored in a Collection.
- Can have `_id` key – works like Primary keys in MySQL.
- Supported Relationships – Embedded (or) References.
- Document storage in BSON (Binary form of JSON).

Understanding the Document Model.

```
var p = {  
  '_id': '3432',  
  'author': DBRef('User', 2),  
  'title': 'Introduction to MongoDB',  
  'body': 'MongoDB is an open sources.. ',  
  'timestamp': Date('01-04-12'),  
  'tags': ['MongoDB', 'NoSQL'],  
  'comments': [{ 'author': DBRef('User', 4),  
                  'date': Date('02-04-12'),  
                  'text': 'Did you see.. ',  
                  'upvotes': 7, ... }  
]  
}  
> db.posts.save(p);
```



Secondary Indexes

Create Index on any field in the document

// 1 means ascending, -1 means descending

```
> db.posts.ensureIndex({'author': 1});
```

//Index Nested Documents


```
> db.posts.ensureIndex('comments.author': 1);
```

// Index on tags

```
> db.posts.ensureIndex({'tags': 1});
```

// Geo-spatial Index

```
> db.posts.ensureIndex({'author.location': '2d'});
```



What about Queries? So

Simple

// find posts which has 'MongoDB' tag.

```
> db.posts.find({tags: 'MongoDB'});
```

// find posts by author's comments.

```
> db.posts.find({'comments.author':  
DBRef('User',2)}).count();
```


// find posts written after 31st March.

```
> db.posts.find({'timestamp': {'gte': Date('31-03-12')}});
```

// find posts written by authors around [22, 42]

```
> db.posts.find({'author.location': {'near':[22, 42]}});
```

\$gt, \$lt, \$gte, \$lte, \$ne, \$all, \$in, \$nin, count, limit, skip, group, etc...



Whatabout Updates? Atomic Operations makes it simple


```
db.posts.update({_id: '3432'},  
{'title': 'Introduction to MongoDB (updated)',  
'text': 'Updated text',  
${addToSet: {'tags': 'webinar'}});
```

\$set, \$unset

\$push, \$pull, \$pop, \$addToSet


\$inc, \$decr, many more...

Where are my joins and transactions? !!!



Some Cool features


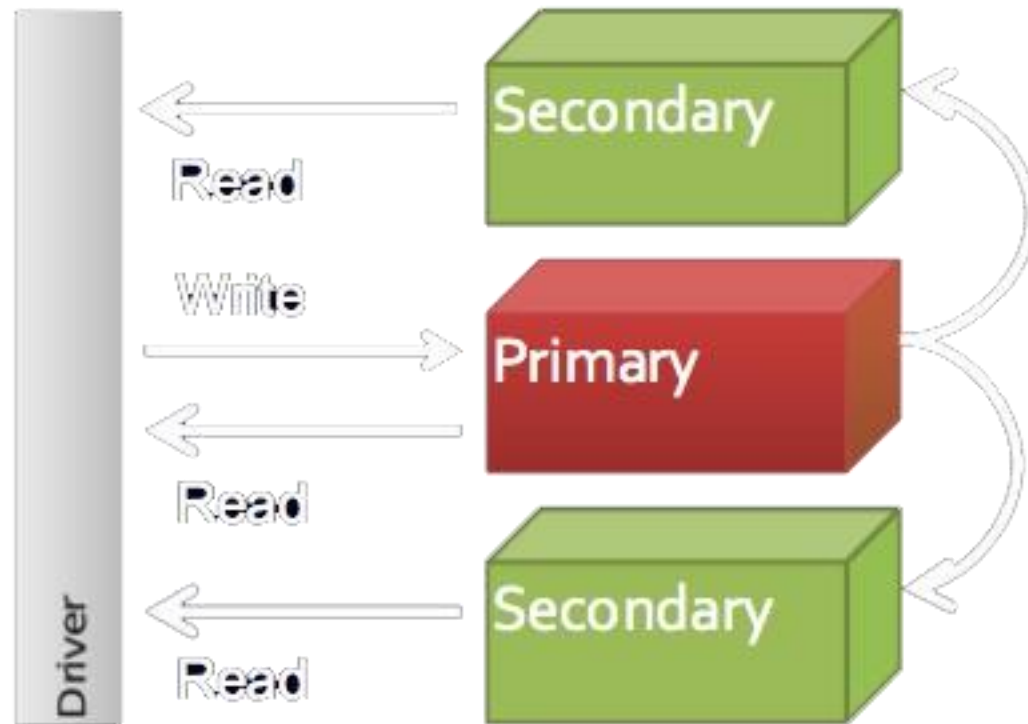
- Geo-spatial Indexes for Geo-spatial queries.
\$near, \$within_distance, Bound queries (circle, box)
- GridFS
Stores Large Binary Files.
- Map/Reduce
GROUP BY in SQL, map/reduce in MongoDB.



Deployment & Scaling



Replica Sets



Sharding

